

PhilanthroBot: A Trust-Centric Conversational Agent for NGO Discovery and Recommendation using Stateful RAG Architecture

Project Review – 1 (Capstone Project)

Presenters:

- 22BIT0013 – Tanish Maheshwari
- 22BIT0100 – Manya Dsouza



The Trust Deficit in Philanthropy

The Problem

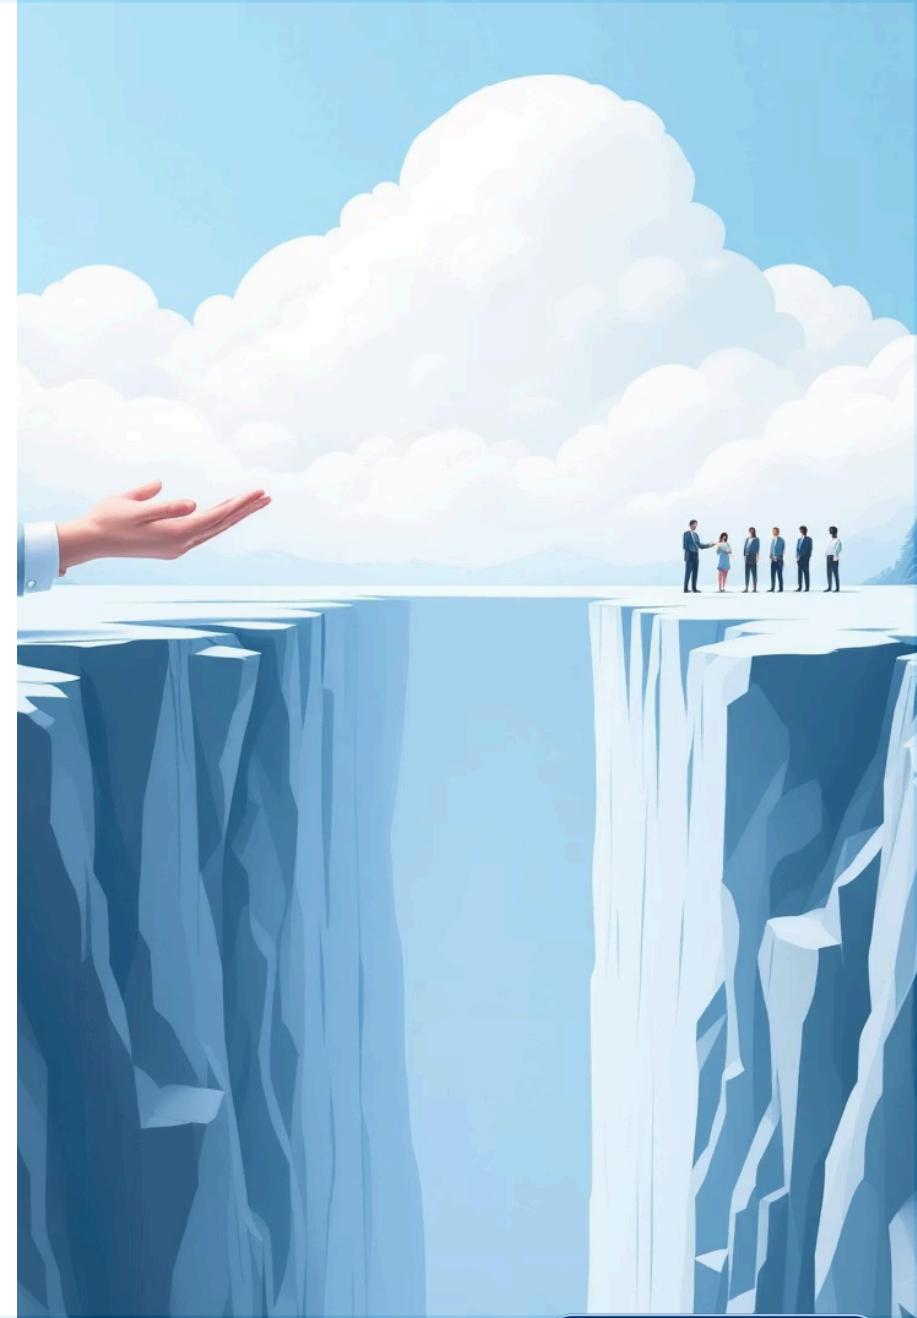
The non-profit sector faces a significant "trust deficit." Donors struggle to vet NGOs due to information asymmetry and lack of transparency.

Context

This isn't a lack of generosity, but a breakdown in the channels connecting donors to NGOs. Donors hesitate without clear accountability on how funds will be used.

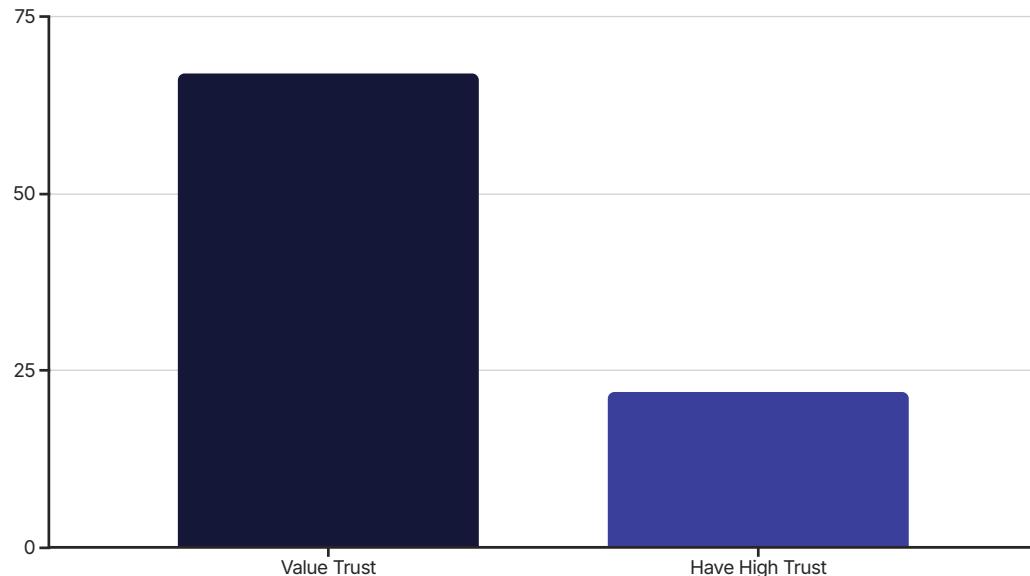
Impact

As a result, many well-intentioned donations never happen. This uncertainty and doubt stall potential contributions, highlighting the gravity of the problem.



Donor Trust – Facts & Figures

Research highlights a significant gap between donors' expectations and their actual trust in charitable organizations, impacting donation decisions.



Donor Concerns

- A major deterrent for **34% of donors** is uncertainty about how their donation will be used.
- A common perception exists that **too much of the donation** might be spent on overhead instead of directly on the cause.

NGO Perspective

NGOs struggle to demonstrate tangible impact and transparency due to **limited resources**, creating a vicious cycle where lack of perceived transparency makes funding scarce.

Gap in Current Solutions



Outdated Platforms

Static NGO websites and donation portals offer basic filters and information, failing to foster genuine trust or deeper understanding.

"Leaky" Donor Funnel

Confusing sites and unclear missions lead to donor drop-off at multiple points, losing potential contributions.

Lack of Personalization

One-size-fits-all platforms treat NGO search like a database query, providing no tailored guidance to build donor confidence.

Traditional search approaches fail to engage donors or address the deeper trust issues needed to inspire meaningful giving.

Proposed Solution – Introducing PhilanthroBot

PhilanthroBot is a trust-centric conversational AI agent designed to revolutionize NGO discovery and recommendation for donors.



Conversational Approach

An AI chatbot guiding donors through natural, interactive dialogue, moving beyond static information browsing.



Adaptive Intelligence

Adapts to donor interests (e.g., environmental causes, specific regions) to tailor NGO recommendations.



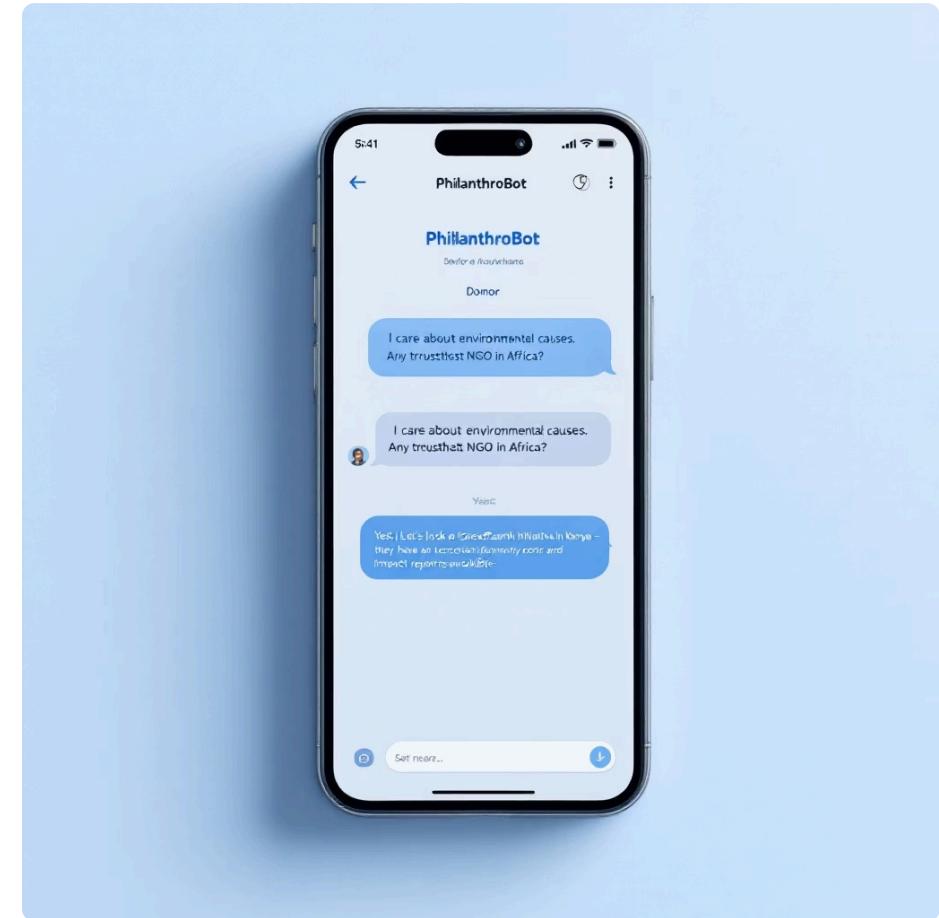
Transparency through Q&A

Delivers real-time, fact-based answers on NGO mission, fund utilization, and impact, fostering confidence.



Empowering Donors

Transforms discovery into a collaborative, trust-building journey, ensuring donors feel heard, informed, and assured.



Project Objectives

PhilanthroBot is designed with clear goals to address the trust deficit and enhance donor engagement.

→ **Ensure Factual Accuracy**

Eliminate AI "hallucinations" using Retrieval-Augmented Generation (RAG) for verified, reliable NGO data.

→ **Personalized Recommendations**

Tailor NGO suggestions based on donor preferences, remembering favored causes and regions.

→ **Trust-Centric Knowledge Base**

Build a specialized NGO database emphasizing transparency, accountability, and impact metrics.

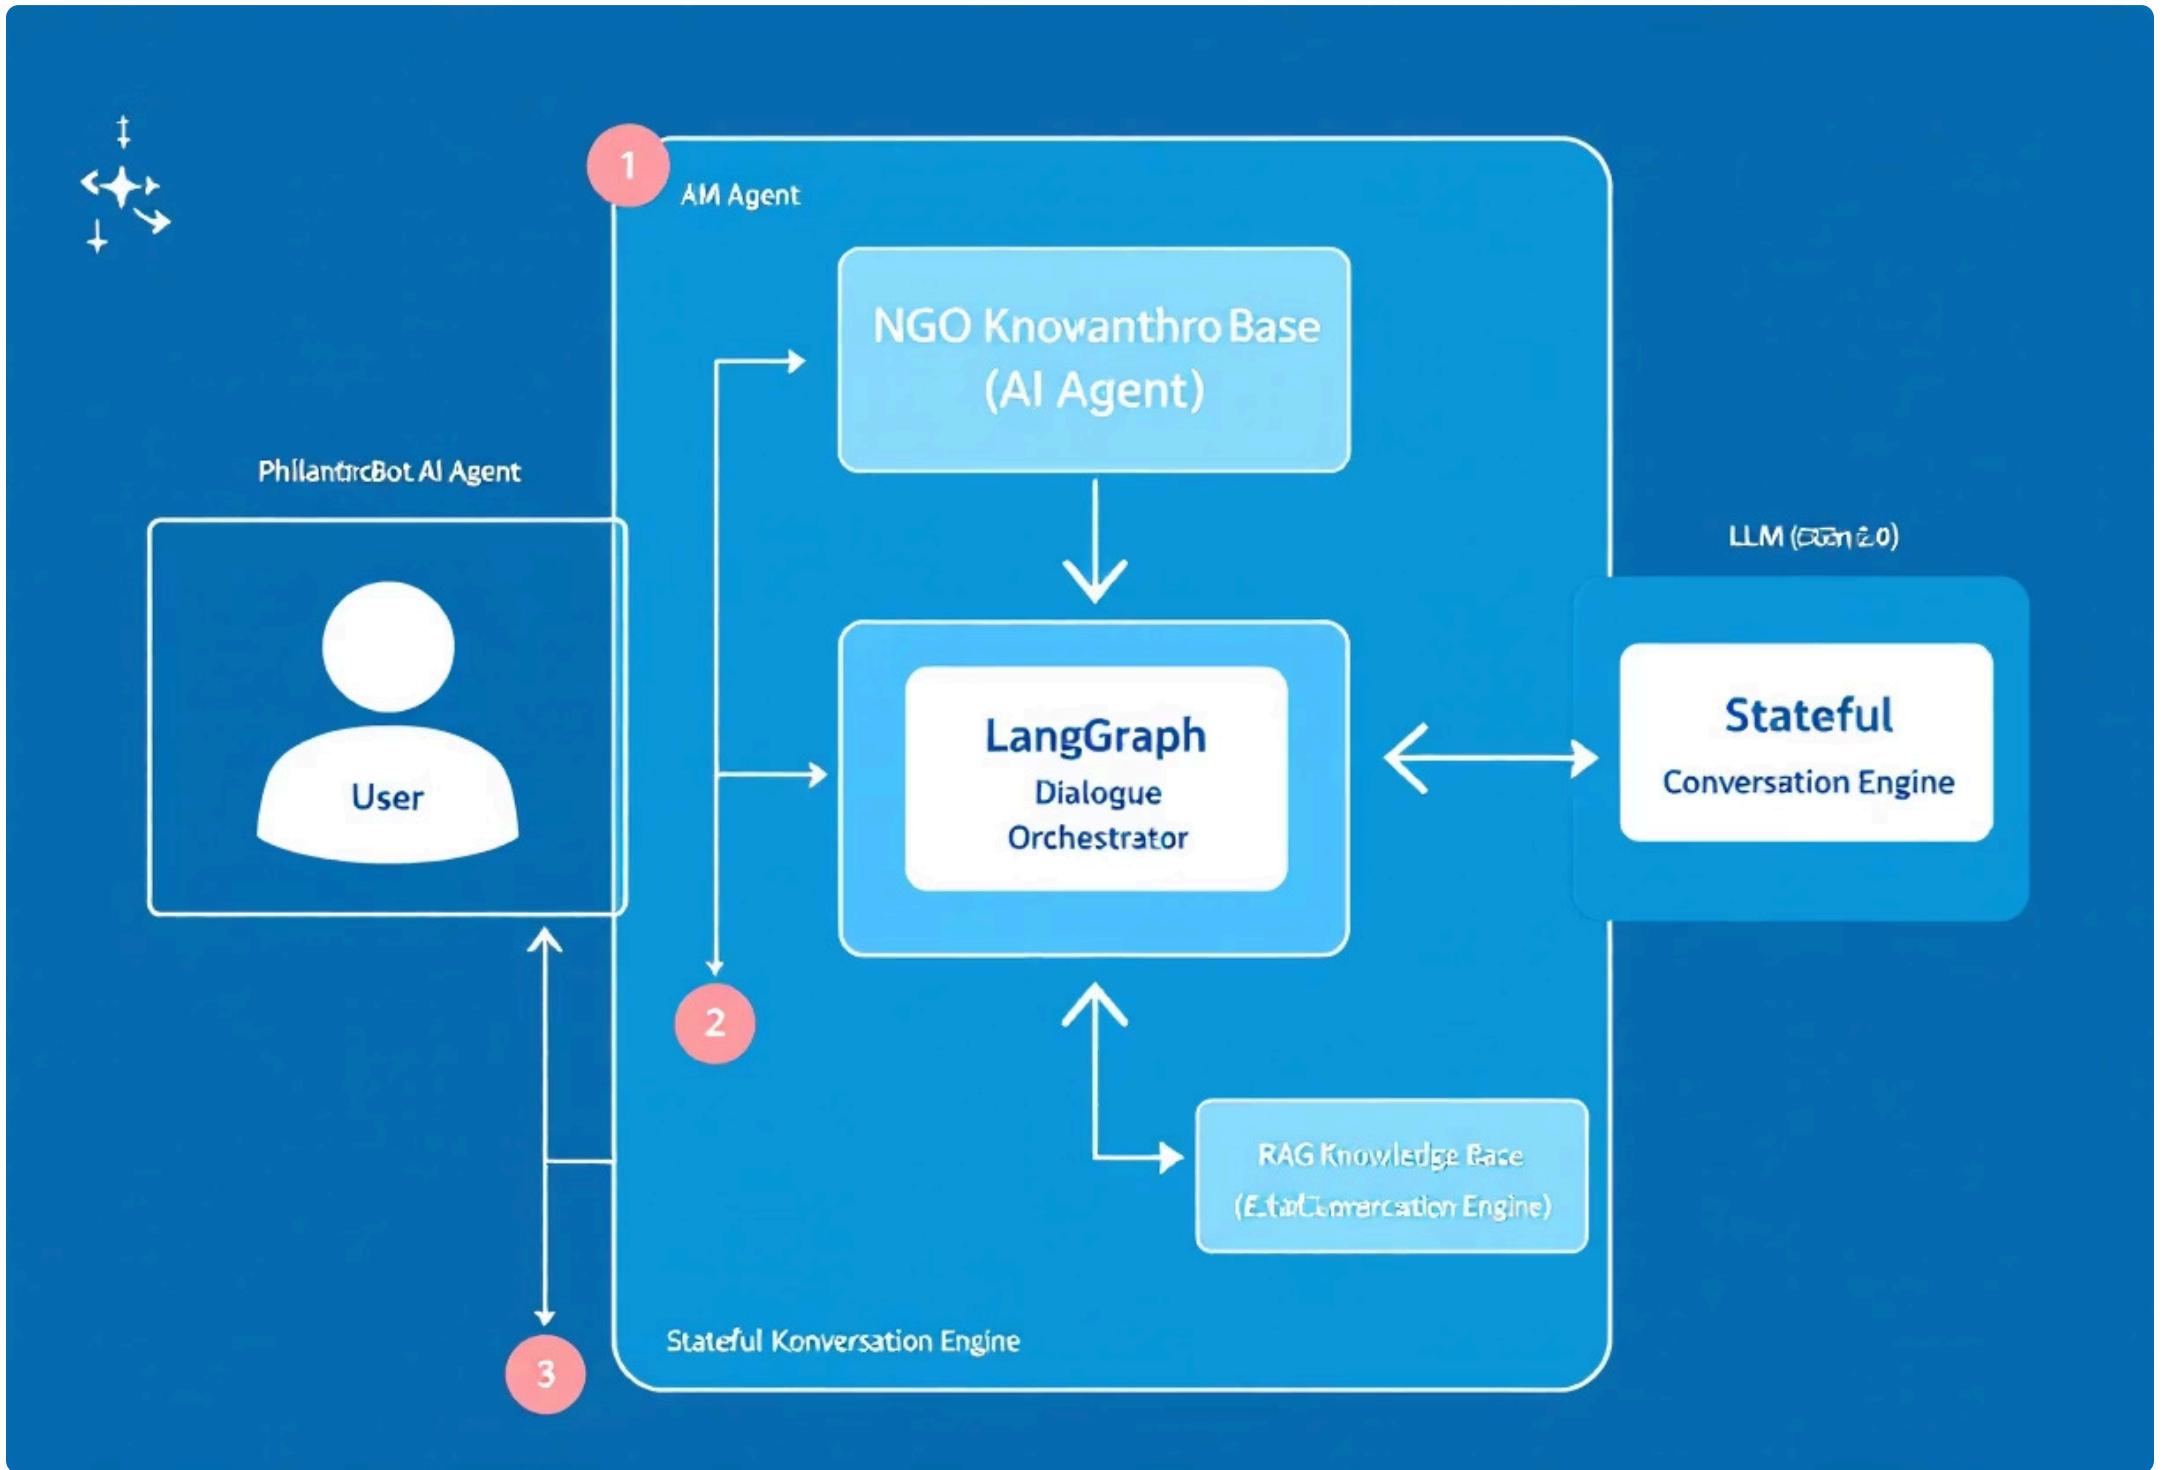
→ **High-Quality User Experience**

Deliver a fast, responsive, and intuitive chatbot interaction leveraging Gemini-2.0-flash.

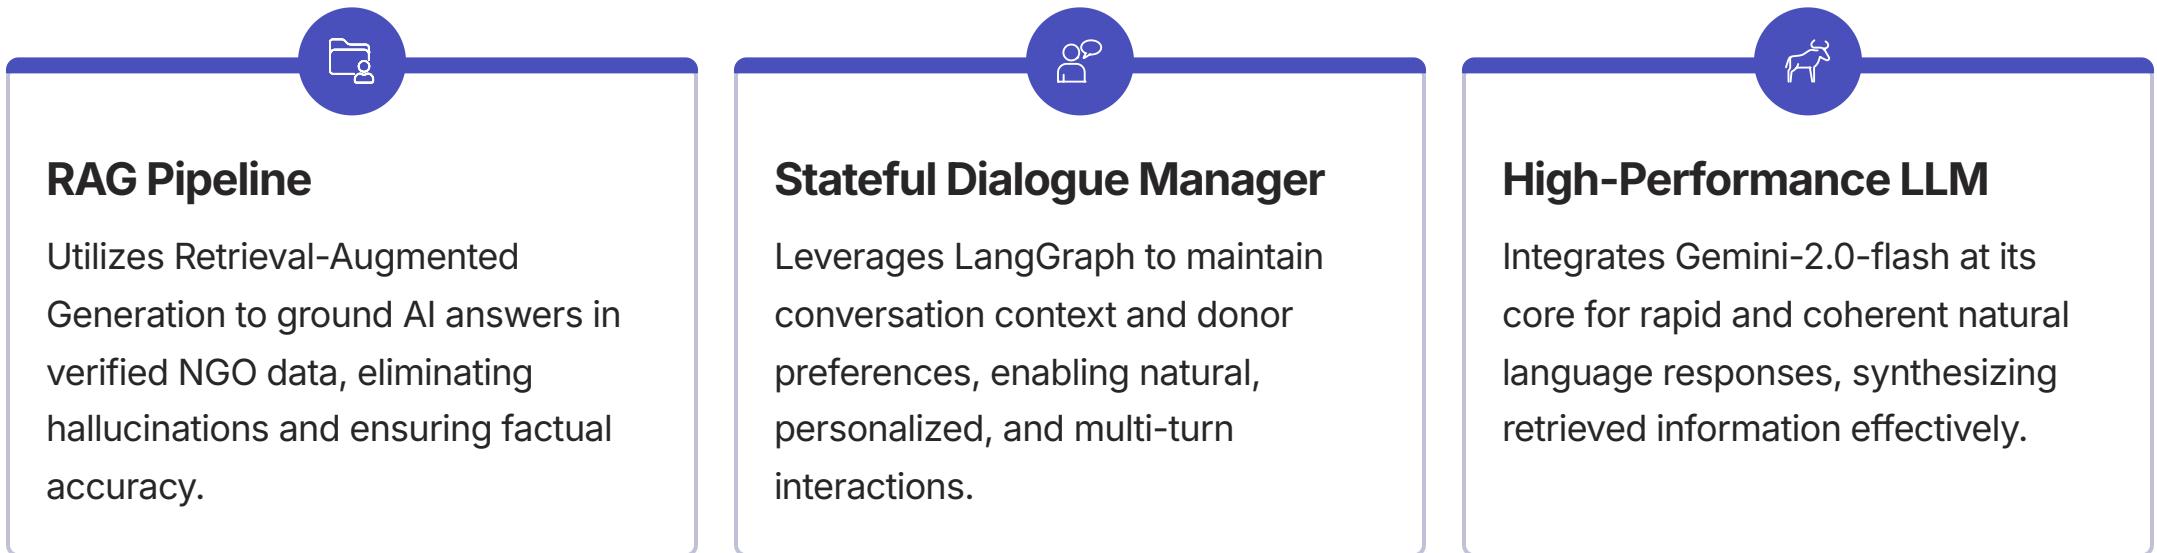


System Architecture Overview

PhilanthroBot is built on a robust, trust-centric architecture, combining advanced AI with a structured knowledge base to deliver accurate and personalized donor recommendations.



Architecture Highlights: Building Trust Through Design

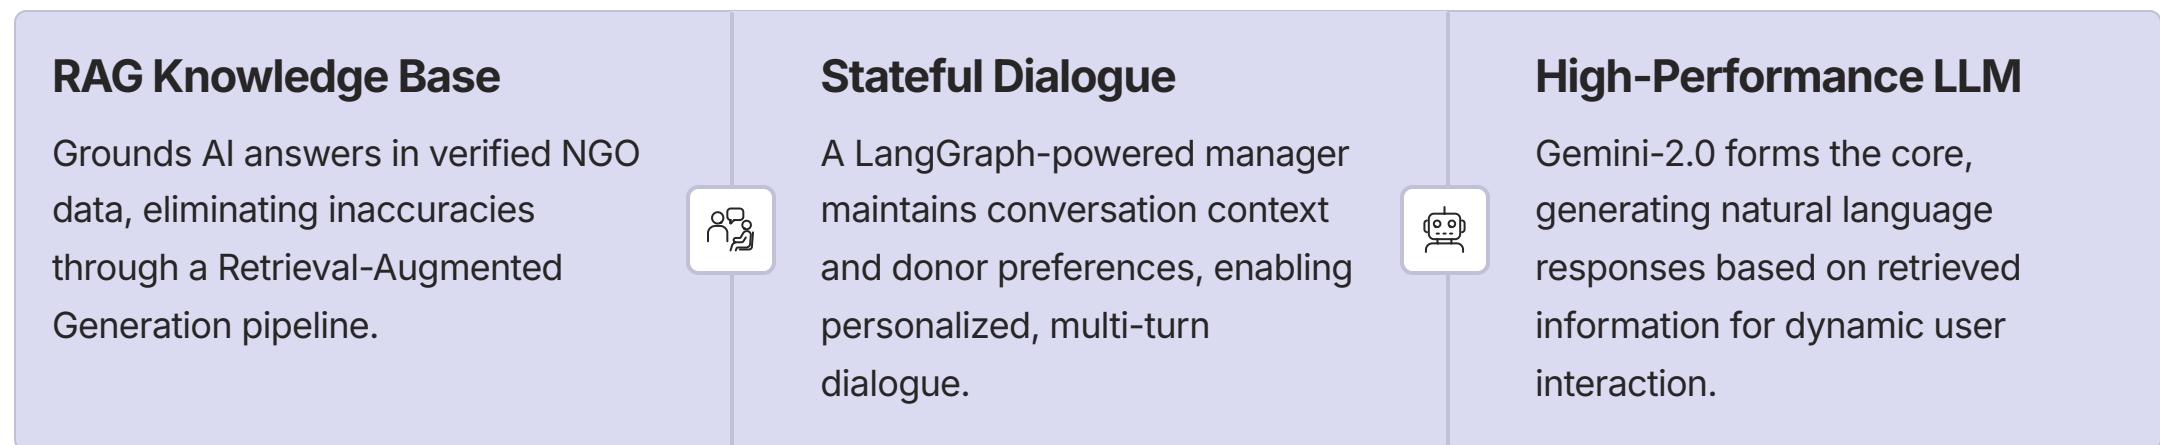


When a donor asks a question, PhilanthroBot first retrieves relevant information from the comprehensive NGO database—considering both the query and the donor's stated preferences. This information then feeds into the language model, which generates a precise and data-backed answer, finally presented to the user. This systematic flow ensures every recommendation is credible and transparent.



System Architecture Overview

PhilanthroBot leverages a robust architecture to ensure trust and personalized interactions, integrating key components for seamless NGO discovery.



When a donor asks a question, PhilanthroBot intelligently retrieves relevant information from the NGO database, then uses the language model to generate a precise, data-backed answer, assuring donors of factual accuracy and transparency.



Phase 1 – Building the Trust Knowledge Base (RAG)

This foundational phase constructs the Retrieval-Augmented Generation (RAG) knowledge base, ensuring PhilanthroBot's factual accuracy and trustworthiness.

01

Curate NGO Profiles

Structured dummy profiles, detailing mission, impact, and financials, are created to simulate real NGO data.

02

Ingest Data

An automated pipeline efficiently loads these diverse NGO documents into the system, preparing them for AI processing.

03

Chunk Text

Documents are broken into small, semantically coherent chunks, allowing for highly relevant information retrieval.

04

Embed & Index

Each text chunk is converted into a numeric vector and stored in a vector database for rapid, accurate similarity searches.

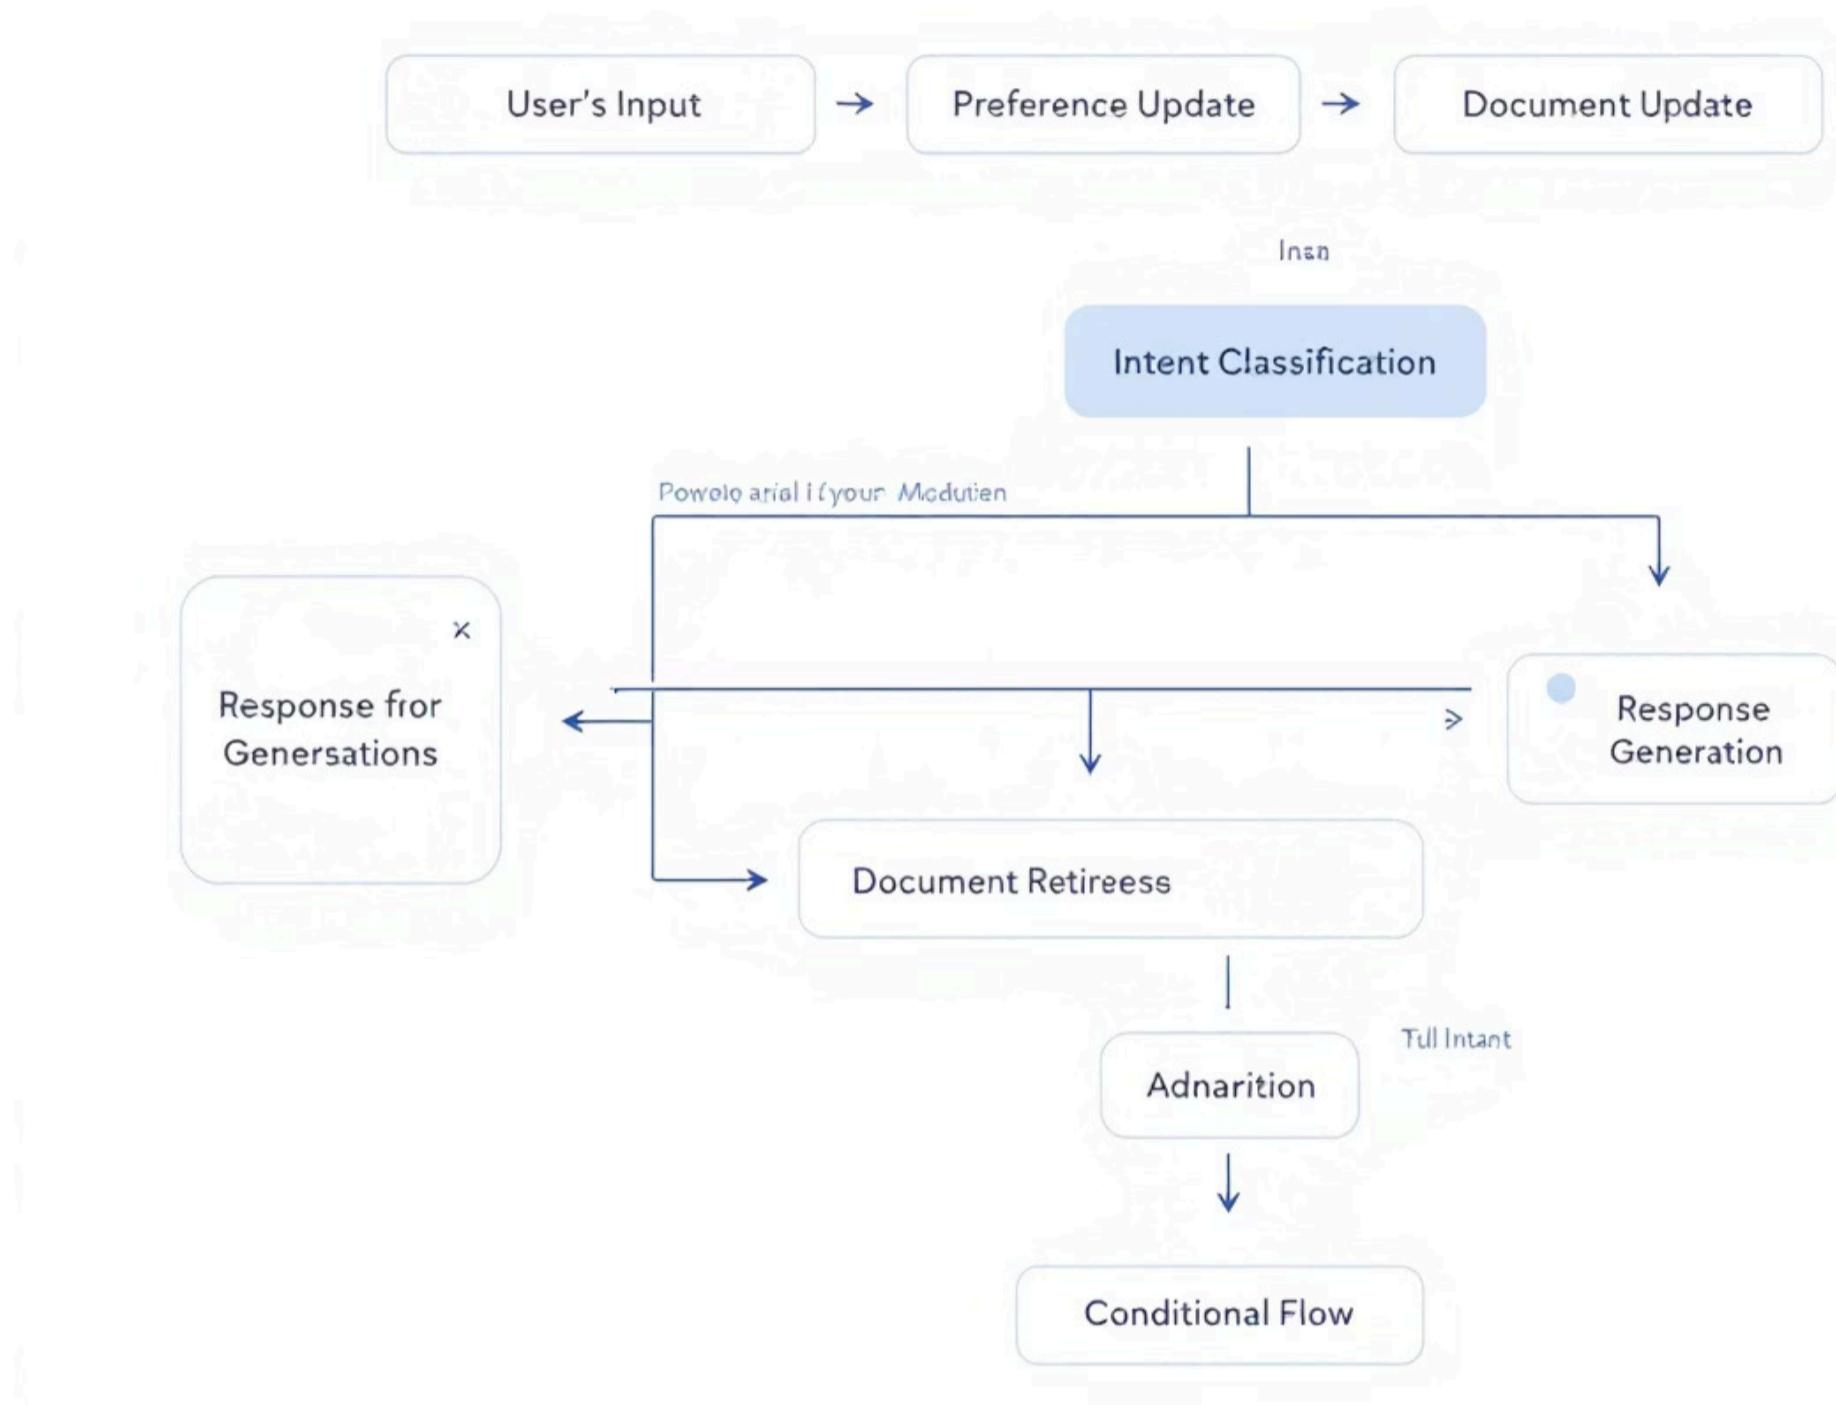
The result is a dynamic, searchable knowledge base, allowing the chatbot to provide real-time, traceable NGO data to build donor trust.

Phase 2 – Developing the Conversational Agent

This phase focuses on building PhilanthroBot's intelligence, enabling it to understand user intent, remember context, and deliver accurate, personalized recommendations through a dynamic conversational flow.

Stateful Dialogue Management

At its core, PhilanthroBot maintains an internal "memory" that persists throughout the conversation. This **persistent state** tracks user preferences (e.g., preferred cause areas, locations, past interactions) and chat history, allowing the bot to avoid repetition, personalize future responses, and create a seamless, natural dialogue experience.



LangGraph Agent Logic: The Brain of PhilanthroBot

The chatbot's decision-making process is orchestrated by a graph of interlinked nodes, implemented using **LangGraph**. This architecture allows for a flexible and intelligent conversation flow, adapting to user input dynamically.

01

Intent Classification

Analyzes user messages to determine their purpose: asking a question, stating a new preference, or following up on a previous topic.

02

Preference Update

If a new preference is detected (e.g., "I want to help children"), the bot's internal state is immediately updated.

03

Document Retrieval

For factual questions, relevant information chunks are efficiently retrieved from the specialized NGO vector database.

04

Response Generation

Utilizes the LLM (Gemini) to formulate coherent, natural language responses, ensuring they are factually grounded by the retrieved information.

05

Conditional Flow

LangGraph's stateful logic connects these nodes, allowing the conversation to branch appropriately based on user input, ensuring adaptability and relevance.

By integrating with Gemini, PhilanthroBot ensures that while its dialogue is natural and engaging, it is always anchored by credible data, building trust with every interaction.

Trust-Centric Data Design

PhilanthroBot's knowledge base isn't just a collection of NGOs; it's meticulously structured to directly address donor concerns about trust and transparency.

Structured Profiles

Each NGO profile contains specific fields critical for donor trust, ensuring that key information is readily accessible for retrieval and verification.

- **Mission Statement:** Core purpose and objectives.
- **Impact & Outcomes:** Measurable results and achievements.
- **Financial Transparency:** Breakdown of fund utilization (e.g., % to programs vs. overhead).
- **Governance & Accreditation:** Details on leadership and third-party certifications.

Example: Trust in Action

For instance, the "**Financial Transparency Summary**" clearly states the percentage of funds allocated to programs versus administrative costs. Similarly, "**Governance & Accreditation**" highlights certifications from watchdog organizations, directly building confidence.

Direct Utility

This intentional design allows PhilanthroBot to provide precise, data-backed answers to donor questions like, "How transparent is this NGO's finances?" or "What verifiable impact have they achieved?" This quick access to verifiable details fosters significant donor confidence.

Ensuring Factual Accuracy – RAG Approach

PhilanthroBot tackles the critical challenge of AI "hallucinations" head-on, ensuring every piece of information delivered is factual and trustworthy.

✖ The Hallucination Problem

Traditional large language models (LLMs) can sometimes generate plausible-sounding but incorrect information. In a donor context, this would be disastrous, as misinformation immediately erodes trust.



Donor Query

A donor asks PhilanthroBot a question about an NGO.



Data Retrieval

Before generating a response, PhilanthroBot retrieves verified data from its comprehensive NGO knowledge base.



Grounded Generation

The AI then uses this specific, verified information as context to generate its answer, preventing guesswork.



Trustworthy Answer

PhilanthroBot provides a precise, data-backed response, often with source attribution.

This Retrieval-Augmented Generation (RAG) approach grounds every response in verified facts, not AI conjecture. By consistently delivering answers based on real, attributable data, PhilanthroBot earns and maintains the donor's trust, transforming discovery into a reliable and transparent experience.

Personalized, Stateful Conversations

PhilanthroBot elevates donor interaction beyond simple FAQs by remembering context and adapting to individual needs, creating a truly personalized experience.



Memory of Context

Recalls prior preferences (e.g., "education projects") for coherent, multi-turn dialogues, making every interaction relevant.



Adaptive Interaction

Refines recommendations and asks clarifying questions based on past input, preventing repetitive queries and streamlining discovery.



Human-Like Personalization

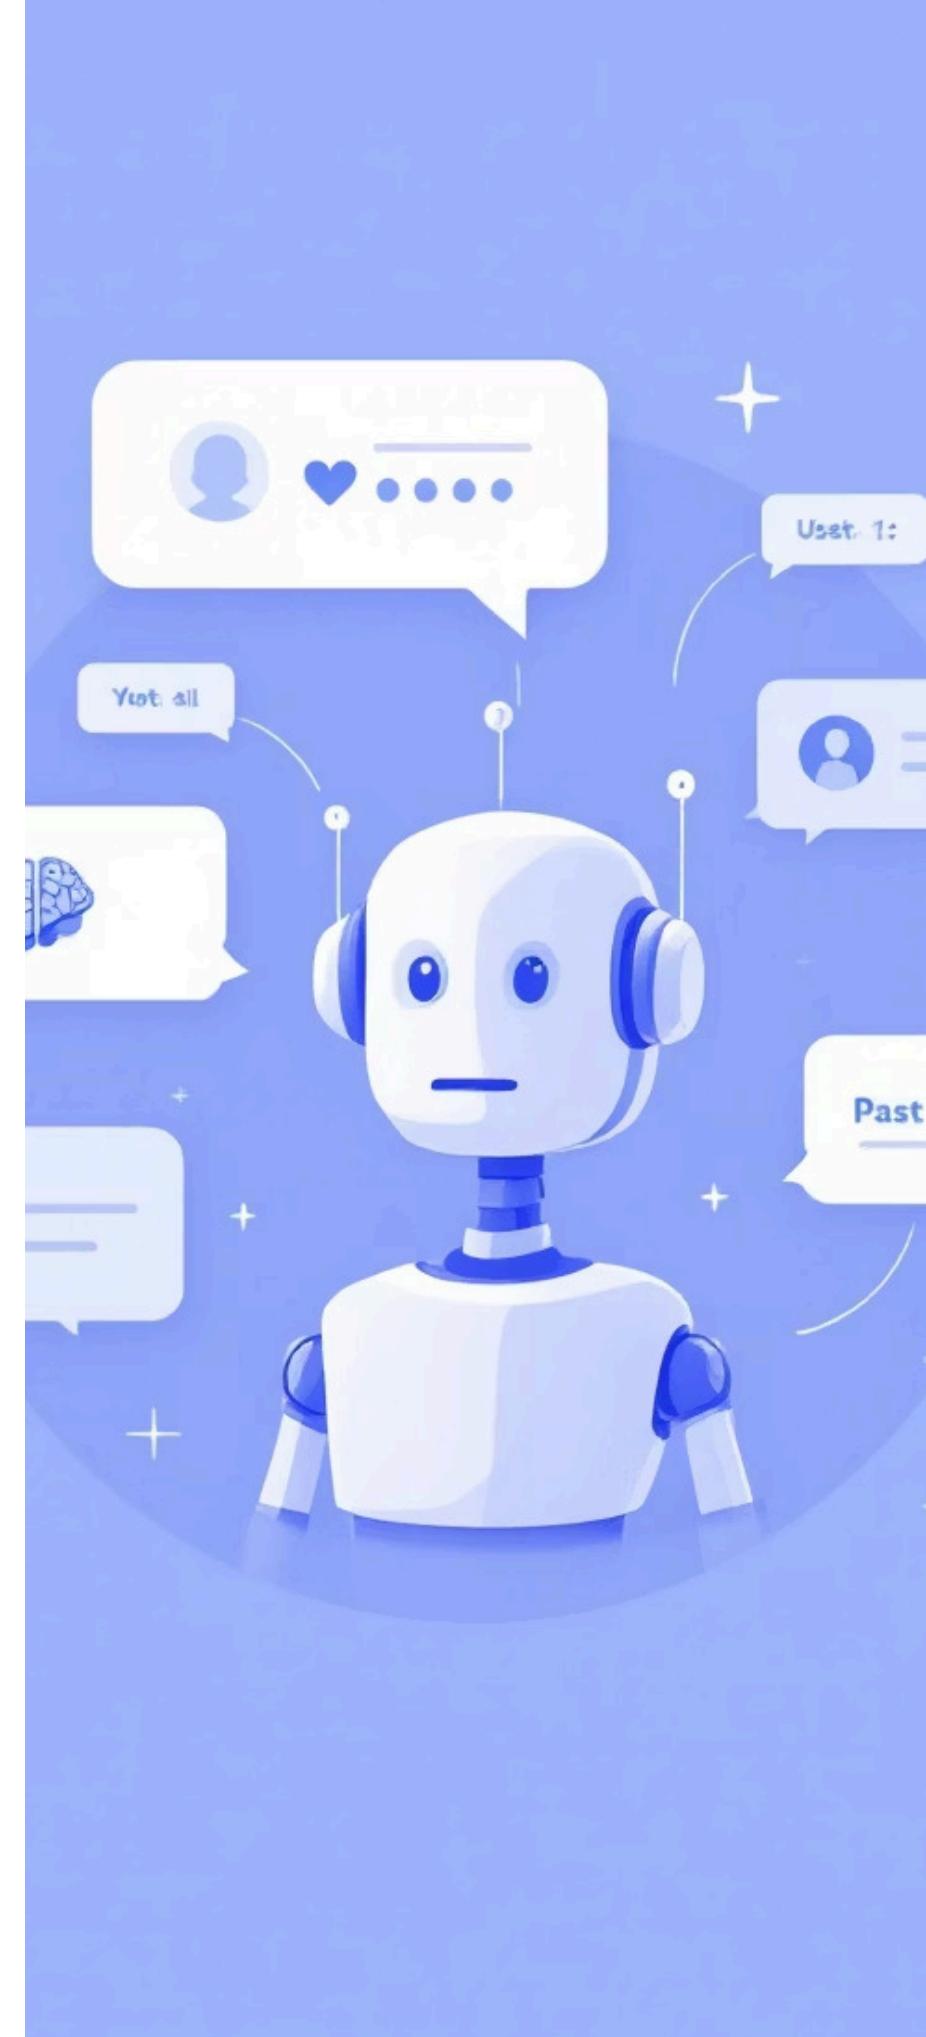
Mimics a knowledgeable human advisor, offering tailored recommendations that feel like trusted advice from a friend.



Trust Through Engagement

Fosters a two-way dialogue, allowing donors to ask follow-up questions and address doubts, deepening engagement and building confidence.

This intelligent, stateful approach ensures that every donor feels heard and understood, transforming the search for NGOs into a deeply engaging and trustworthy journey.



Tech Stack & Implementation Details

PhilanthroBot is built on a robust foundation of modern AI and data technologies, carefully selected for their efficiency, scalability, and ability to ensure a trustworthy user experience.



Gemini-2.0-Flash (LLM)

Chosen for its **speed** and **large context window**, enabling real-time responses and handling extensive NGO information snippets effectively.



ChromaDB & Embeddings

This vector database stores NGO data chunks as vectors, facilitating **quick semantic search** for relevant information. It uses **Google Generative AI Embeddings**.



LangChain & LangGraph

These powerful frameworks **accelerate development** by integrating RAG and structuring the conversational flow, managing state and routing logic efficiently.



Deployment & Interface

Starting with a web chat UI, we're exploring deployment on **WhatsApp** or similar messaging platforms for convenient, on-the-go donor interaction.

This strategic technology stack ensures PhilanthroBot delivers on its promise of accurate, personalized, and trust-centric NGO discovery.

Current Status & Progress

We've made significant strides, laying a solid foundation and pushing key development phases forward. Here's an overview of our achievements and ongoing work.

1

Literature Review & Planning

Status: Completed. Thorough analysis of the trust deficit and existing solutions. Defined a clear problem statement, rationale, and a well-structured project plan and design documents.

2

Knowledge Base Prototype

Status: 80% Completed. We've successfully implemented the Phase 1 pipeline, ingesting, chunking, and indexing dummy NGO profiles. Test queries confirm data retrieval.

3

Conversational Agent Development

Status: 40% Completed. Basic conversation flow logic is set up, including state structure and initial functions (intent classification). A command-line chatbot can greet and perform basic fact lookups.

4

Early Demo Results

Status: Proof-of-Concept Achieved. Successfully demonstrated a query like "Tell me about NGO X's mission," with the system fetching data and the LLM framing a proper response. This validates core component integration.

Next Steps: Integrate full LLM responses, refine dialogue management, and expand the knowledge base.

Our team is coordinating effectively, with dedicated members focusing on the knowledge base and chatbot logic, ensuring continuous progress and alignment with project goals.

Project Timeline & Milestones

Our journey from inception to completion, highlighting key phases and deliverables.



Next Immediate Goal: By end of September, a working end-to-end demo of PhilanthroBot with a small set of NGOs.

Challenges & Risk Mitigation

Anticipating potential hurdles and proactively developing strategies to ensure project success and maintain donor trust.

Data Quality & Validity

Challenge: The system's answers are only as good as its data. Ensuring NGO profiles are accurate, up-to-date, and unbiased is critical.

Mitigation: We plan to exclusively use reputable sources for data (e.g., verified NGO reports) and implement periodic updates. A verification step for critical data points will be incorporated before they are added to the knowledge base.

Technical Integration Complexity

Challenge: Combining diverse components like the LLM, vector database, and LangGraph can lead to complex integration bugs or mismatches, potentially causing incorrect or missing answers.

Mitigation: We're employing an incremental testing approach: first ensuring robust data retrieval, then gradually integrating the LLM, and finally the full conversational dialogue. This modular strategy allows us to catch and address issues early.

LLM Behavior & Limitations

Challenge: Even with RAG, the LLM might occasionally produce irrelevant or overly verbose responses, impacting user experience.

Mitigation: We will continuously refine our prompts and implement rules within the conversation logic (e.g., truncating responses or re-querying if an answer seems off). Exploring alternative models or fine-tuning for our specific domain will be considered if needed.

User Adoption & Trust

Challenge: Despite our focus on trust, users unfamiliar with AI may be skeptical of a chatbot for donation discovery.

Mitigation: Our UX design prioritizes transparency, allowing the bot to display sources on request. We will also provide user education during onboarding and focus on crafting smooth, helpful initial interactions to quickly build user confidence.

Resource Constraints

Challenge: Utilizing a large model like Gemini-2.0-Flash can be computationally intensive, potentially impacting performance or incurring high costs.

Mitigation: We are leveraging efficient cloud resources for processing and will judiciously scope the knowledge base for initial demos to manage computational demands. Future scaling will be optimized through strategic resource allocation.

By systematically addressing these challenges, we are building a resilient and trustworthy PhilanthroBot.

Expected Outcomes & Impact

Our vision for success extends beyond development, encompassing both tangible deliverables and a profound impact on the philanthropic landscape.



Functional Prototype

Delivery of a complete, working PhilanthroBot with a robust back-end (knowledge base and conversational agent) and an intuitive user interface for demonstrations.



Enhanced Donor Trust

Qualitatively improve user confidence compared to standard directories, reducing decision paralysis and increasing satisfaction through interactive, factual Q&A.



Technical Showcase

Illustrate the effective integration of LangChain, LangGraph, and LLMs, offering a potential case study or academic paper on AI-driven trust improvement.



Donor Query Guidance

Demonstrate the bot's ability to handle diverse donor queries, from cause and location searches to credibility checks and NGO comparisons, guiding users to vetted recommendations.



Scalable Blueprint

Establish a foundational framework for trust-centric AI assistants in the non-profit sector, extensible to real NGO databases and other trust-sensitive domains.

These outcomes will solidify PhilanthroBot's potential to bridge the trust gap and redefine donor engagement in philanthropy.

Conclusion

PhilanthroBot was conceived to address the critical trust deficit in charitable giving. By leveraging grounded AI (RAG) for factual accuracy and stateful personalization, it transforms the donor experience into a guided, trustworthy journey.

This innovation enhances transparency and engagement, fostering deeper connections between donors and NGOs, ultimately leading to increased support for vital causes.

PhilanthroBot – Bringing Trust into the Heart of Giving.

Questions? Thank You!

We welcome any questions or feedback on PhilanthroBot and our journey to build trust in philanthropy.